

Unsupervised Analysis: Dimension Reduction

Why Dimension Reduction?

For Big-Data:

- Data visualization becomes very difficult! (Cannot draw 2D scatterplots between all pairs of features).
- Big-Data often has a high degrees of redundancy. (i.e. correlation among features).
- Many features may be uninformative for the particular problem under study (noise features).
- Dimension reduction ideally allows us retain information on most important features of the data, while reducing noise and simplifying visualization & analysis.

What is Dimension Reduction?

- Map the data into a new low-dimensional space where important characteristics of the data are preserved.
- The new space often gives a (linear or non-linear) transformation of the original data.
- Visualization and analysis (clustering/prediction/...) is then performed in the new space.
- In many cases, (especially for non-linear transformations) interpretation becomes difficult.

Principal Components Analysis (PCA)

PCA

Set-up:

- Data matrix: $\mathbf{X}_{n \times p}$, n observations and p features.

Idea:


- Not all p features are needed (much redundant info).
- Find low-dimensional representations that capture most of the variation in the data.

Uses:

- Ubiquitously used - Dimension reduction, data visualization, pattern recognition, exploratory analysis, etc.
- Best linear dimension reduction possible.

PCA - Main Idea

Question: What is a good 1D representation of the data?



PCA - Main Idea

Some Possibilities:


- Use one of the variables (e.g. x_1).
- Better idea: use a linear combination of the variables (i.e. a weighted average).

$$z_1 = v_1 x_1 + v_2 x_2 = \mathbf{v}^T \mathbf{X}$$

How to choose the weights (v_1 and v_2)?

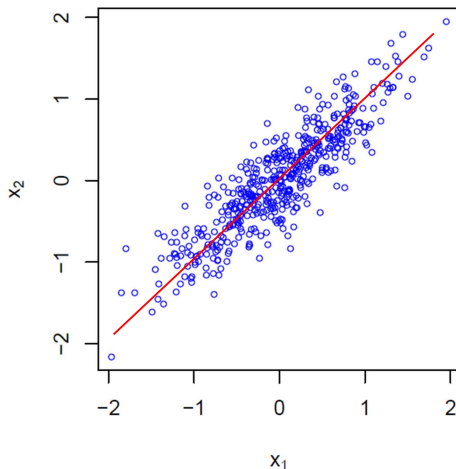
PCA - Main Idea

Many possibilities, but which one is a good choice?



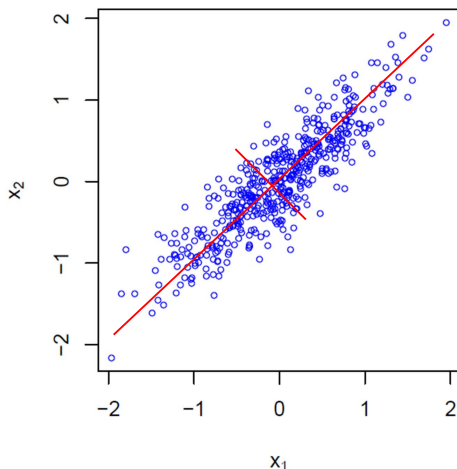
PCA - Main Idea

Find line that maximizes the variance of the data projected onto the line:




PCA - Main Idea

Subsequent components orthogonal (perpendicular).





PCA - Main Idea



- PCA minimizes orthogonal projection onto line: $Z = v_1x_1 + v_2x_2$.
- Slope of line = v_2/v_1 (if features centered).
- Note: Not same as OLS which minimizes projection of y onto x !

PCA - Main Idea

3D Projection onto a Hyperplane:



PCA - Criterion

PCA Criterion - PC 1 (Population):

$$\underset{\mathbf{v}}{\text{maximize}} \quad \text{Var}(\mathbf{X}\mathbf{v}) \quad \text{subject to } \|\mathbf{v}\|_2 = 1$$

$$\underset{\mathbf{v}}{\text{maximize}} \quad \mathbf{v}^T \text{Var}(\mathbf{X}) \mathbf{v} \quad \text{subject to } \|\mathbf{v}\|_2 = 1$$

$$\underset{\mathbf{v}}{\text{maximize}} \quad \mathbf{v}^T \mathbf{\Sigma} \mathbf{v} \quad \text{subject to } \|\mathbf{v}\|_2 = 1$$

where $\mathbf{\Sigma} = \text{Cov}(\mathbf{X})$.

- Finds linear combination of features that maximizes the variance.

PCA - Criterion

PCA Criterion - PC k (Population):

$$\underset{\mathbf{v}_k}{\text{maximize}} \quad \mathbf{v}_k^T \boldsymbol{\Sigma} \mathbf{v}_k \quad \text{subject to} \quad \|\mathbf{v}_k\|_2 = 1 \text{ \& } \mathbf{v}_k^T \mathbf{v}_j = 0 \quad \forall j < k.$$

- Subsequent linear combinations are orthogonal to previous combinations.
- **Uncorrelated.**

PCA - Criterion

PCA Criterion - Sample Version:

$$\underset{\mathbf{v}_1, \dots, \mathbf{v}_K}{\text{maximize}} \quad \mathbf{v}_k^T \mathbf{X}^T \mathbf{X} \mathbf{v}_k \quad \text{subject to } \|\mathbf{v}_k\|_2 = 1 \text{ \& } \mathbf{v}_k^T \mathbf{v}_j = 0 \quad \forall j < k.$$

Replaces Σ with estimate $\mathbf{X}^T \mathbf{X} / n$.

Solution: Eigenvalue decomposition of $\mathbf{X}^T \mathbf{X}$. (`eigen()` in R)

PCA - Criterion

Equivalent PCA Criterion:

$$\begin{aligned} & \underset{\mathbf{u}_1, \dots, \mathbf{u}_K, \mathbf{v}_1, \dots, \mathbf{v}_K}{\text{maximize}} && \mathbf{u}_k^T \mathbf{X} \mathbf{v}_k \quad \text{subject to } \|\mathbf{v}_k\|_2 = 1 \ \& \ \mathbf{v}_k^T \mathbf{v}_j = 0 \ \forall j < k. \\ & && \|\mathbf{u}_k\|_2 = 1 \ \& \ \mathbf{u}_k^T \mathbf{u}_j = 0 \ \forall j < k. \end{aligned}$$

- Finds left and right projection that maximize variance.

Solution: Singular Value Decomposition (SVD) of \mathbf{X} . (`svd()` in R)

PCA - Parts of the Solution

$$\text{SVD: } \mathbf{X}_{n \times p} = \mathbf{U}_{n \times n} \mathbf{D}_{n \times p} \mathbf{V}_{p \times p}^T$$

- Singular vectors: (left) \mathbf{U} and (right) \mathbf{V} .
 - ▶ Orthonormal - $\mathbf{U}^T \mathbf{U} = \mathbf{I}$ and $\mathbf{V}^T \mathbf{V} = \mathbf{I}$.
- Singular values: Diagonals of \mathbf{D} .
 - ▶ $d_1 \geq d_2 \geq \dots \geq d_r$ where $r = \text{rank}(\mathbf{X})$.

SVD Solution to PCA:

- **PCs:** $\mathbf{Z} = \mathbf{X}\mathbf{V}$ or $\mathbf{Z} = \mathbf{U}\mathbf{D}$. (\mathbf{U} are un-scaled PCs).
 - ▶ $\mathbf{z}_k = \mathbf{X}\mathbf{v}_k$ - k^{th} PC.
 - ▶ $\mathbf{z}_1 \dots \mathbf{z}_K$ gives best K -dimensional projection of the data.
- **PC Loadings:** \mathbf{V} .
 - ▶ \mathbf{v}_k - k^{th} PC loading (feature weights).

PCA - Properties


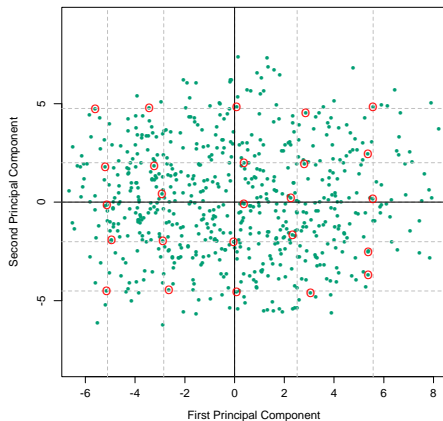
- Unique.
 - ▶ \mathbf{U} and \mathbf{V} unique up to a sign change.
 - ▶ \mathbf{D} unique.
- Global Solution.

PCA - Pattern Recognition

- \mathbf{u}_1 - first column of \mathbf{U} encodes first major pattern in observation space.
- \mathbf{v}_1 - first column of \mathbf{V} encodes the associated first pattern in feature space.
- d_1 gives strength of first pattern.
- Subsequent patterns are **uncorrelated** to first pattern (i.e. orthogonal).
- $\mathbf{X} \approx \sum_{k=1}^K d_k \mathbf{u}_k \mathbf{v}_k^T$ - data is comprised of a series of patterns.

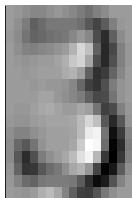
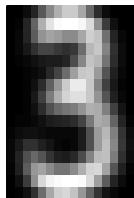
PCA - Pattern Recognition

Patterns in observation space:



PCA - Pattern Recognition

Patterns in feature space:



PCA - Data Visualization

PC Scatterplots:

- Problem: Can't visualize
- Solution: Plot \mathbf{u}_1 vs. \mathbf{u}_2 and so forth.
- Advantages:
 - ▶ Dramatically reduces number of 2D scatterplots to visualize.
 - ▶ Focuses on patterns with most variance.

PC Loadings Plots:

- Scatterplots of \mathbf{v}_1 vs. \mathbf{v}_2 .
- Visualizations of \mathbf{v}_k .



Biplot:

- Scatterplot of PC 1 vs. PC 2 with loadings of \mathbf{v}_1 vs. \mathbf{v}_2 overlaid.

(See demo examples.)

PCA - Data Visualization

Scatterplots:



- Plotting Scatterplot PCs roughly equivalent to rotating axes of original plot.

PCA - Dimension Reduction

Best low-rank approximation to the data:

$$\underset{\tilde{\mathbf{X}}}{\text{minimize}} \quad \|\mathbf{X} - \tilde{\mathbf{X}}\|_F^2 \quad \text{subject to } \text{rank}(\tilde{\mathbf{X}}) = K$$

Solution: $\tilde{\mathbf{X}} = \sum_{k=1}^K d_k \mathbf{u}_k \mathbf{v}_k^T$ - SVD / PCA solution!

- PCA also finds best data compression to minimize reconstruction error.
- PCA yields best linear dimension reduction possible!

PCA - Dimension Reduction

How much variance is explained? (i.e. extent of dimension reduction)

- Variance explained by k^{th} PC:

$$d_k^2 = \mathbf{v}_k^T \mathbf{X}^T \mathbf{X} \mathbf{v}_k.$$

- Total variance of data:

$$\sum_{k=1}^n d_k^2.$$

- Proportion of variance explained by k^{th} PC:

$$d_k^2 / \sum_{k=1}^n d_k^2.$$



- Cumulative variance explained by first r PCs:

$$\sum_{k=1}^r d_k^2 / \sum_{k=1}^n d_k^2.$$

(Extent of dimension reduction achieved by first r PC projections.)

PCA - Dimension Reduction

Screeplot:



PCA - Dimension Reduction

How to choose K ?

- Elbow in screeplot.
- Take K that explains at least 90% (95%, 99%, etc.) variance.
- More sophisticated:
 - ▶ Cross-Validation done internally.
 - ▶ Validation via matrix completion.
 - ▶ Nuclear norm penalties.

PCA - Center and Scale?

- Typically, one should center features (i.e. columns of \mathbf{X}).
 - ▶ Maximizing variance interpretation (assumes multivariate Gaussian model).
- Scaling changes PCA solution.
 - ▶ Features with large scale contribute more to variance, have large PC loadings.


General Suggestions:

- Scale if features measured differently. (Example - US college data).
- Don't scale if features measured in same way & scale has meaning. (Example - gene expression data).


PCA - Applications

“EigenImages” or “EigenFaces”

PCA



PCA - Applications



PCA - Summary

Strengths:

- Best linear dimension reduction.
- Ordered / orthogonal components.
- Unique, global solution.
- Others?

Weaknesses:

- Non-linear patterns.
- Ultra-high-dimensional settings ($p \gg n$)
- Others?

Sparse PCA

Motivation:

- When $p \gg n$, many features irrelevant.
- PCA can perform poorly.

Idea:

- Sparsity in \mathbf{V} : zero out irrelevant features from PC loadings.
- Advantage: Find important features that contribute to major patterns in the data.

How?

- Typically, optimize PCA criterion with sparsity-encouraging penalty of \mathbf{V} .
- Many methods - active area of research!

In R: SPC in PMA package.

Functional PCA

Motivation:

- Times series, ordered data, spatial data.

Idea:

- Want PC loadings to be smooth (vary continuously) over time or space.
- Advantage: Improve interpretation.

How?

- Typically, optimize PCA criterion with a penalty that encourages smoothness of \mathbf{V} over time or space.
- Many methods for both functional data (data in the form of curves) and discretely-sampled functional data (e.g. discrete time points or specific locations).

In R: package `fpca`.

Kernel PCA

Motivation:

- Non-linear patterns.

Idea:

- Embed inner product distances $(x_i^T x_{i'})$ in a higher-dimensional “kernel” space, $k(x_i, x_{i'})$.
- Kernel examples:
 - ▶ Radial: $k(x_i, x_{i'}) = e^{-\|x_i - x_{i'}\|_2^2 / 2\sigma^2}$.
 - ▶ Polynomial: $k(x_i, x_{i'}) = (cx_i^T x_{i'} + 1)^d$.
- Kernel Matrix: $\mathbf{K}_{n \times n} : \mathbf{K}_{ii'} = k(x_i, x_{i'})$.
 - ▶ Idea: \mathbf{K} a non-linear distance matrix.
- Find major non-linear patterns by performing PCA on \mathbf{K} :

$$\mathbf{K} = \mathbf{U} \mathbf{D}^2 \mathbf{U}^T$$

Supervised Dimension Reduction

Partial Least Squares:

- Best dimension reduction of cross-covariance between \mathbf{X} and \mathbf{Y} such that factors are orthogonal to \mathbf{X} .

Canonical Correlations Analysis:

- Best dimension reduction of cross-covariance between \mathbf{X} and \mathbf{Y} such that bi-projection is orthogonal to \mathbf{X} or \mathbf{Y} .

Linear Discriminant Analysis (classification):

- Best dimension reduction of between class covariance matrix relative to within class covariance.


Non-Negative Matrix Factorization (NMF)

NMF

Idea: $\mathbf{X}_{n \times p} \approx \mathbf{W}_{n \times K} \mathbf{H}_{K \times p} = \sum_{k=1}^K \mathbf{W}_{:,k} \mathbf{H}_{k,:}$ with $K \ll p$.

- $\mathbf{X} \geq 0$ - non-negative data matrix.
- $\mathbf{W} \geq 0$ - non-negative observation factors; often sparse (Basis Factors).
 - ▶ $\mathbf{W}_{:,k} \geq 0$ - k^{th} observation factor.
- $\mathbf{H}_{kj} \geq 0$ - non-negative feature factors; often sparse (Mixture Factors).
 - ▶ $\mathbf{H}_{k,:} \geq 0$ - mixture of features that comprise the k^{th} factor.

Like PCA except finds patterns with same direction of correlation.



NMF Interpretation

Topic Modeling:

- \mathbf{X} a matrix of news articles (rows) by words (columns) whose entries are word counts.
 - ▶ $\mathbf{X} \approx \sum_{k=1}^K \mathbf{W}_{:,k} \mathbf{H}_{k,:}$ - sum of topics.
 - ▶ $\mathbf{X}_{ij} = \mathbf{W}_{i,:}^T \mathbf{H}_{:,j} = \sum_{k=1}^K \mathbf{W}_{ik} \mathbf{H}_{kj}$.
- Topic k : Outer-product of k^{th} column of \mathbf{W} ($\mathbf{W}_{:,k}$) and k^{th} row of \mathbf{H} ($\mathbf{H}_{k,:}$).
 - ▶ E.g. Gay marriage.
- $\mathbf{H}_{k,:}$ non-zeros- words contributing to topic k .
 - ▶ E.g. marriage, gay, Supreme, Court, district, equal, etc.
- $\mathbf{W}_{:,k}$ non-zeros - news articles belonging to topic k .
 - ▶ E.g. "North Carolina Allows Officials to Refuse to Perform Gay Marriages" (*New York Times*).

NMF Criterion - Continuous Data

$$\begin{aligned} & \underset{\mathbf{W}, \mathbf{H}}{\text{minimize}} \quad \|\mathbf{X} - \mathbf{W} \mathbf{H}\|_F^2 \\ & \text{subject to} \quad \mathbf{W}_{ik} \geq 0 \ \& \ \mathbf{H}_{kj} \geq 0 \end{aligned}$$

(PCA criterion except with non-negativity constraints.)

Algorithm Updates: (Alternating Non-negative Least Squares)

$$\begin{aligned} \hat{\mathbf{W}} &= \left(\mathbf{X} \mathbf{H}^T (\mathbf{H}^T \mathbf{H})^{-1} \right)_+ \\ \hat{\mathbf{H}} &= \left((\mathbf{W}^T \mathbf{W})^{-1} \mathbf{W}^T \mathbf{X} \right)_+ \end{aligned}$$

Local Solution.

NMF Criterion - Count Data

$$\begin{aligned} & \underset{\mathbf{W}, \mathbf{H}}{\text{minimize}} && \sum_{i=1}^n \sum_{j=1}^p [\mathbf{x}_{ij} \log(\mathbf{w}_i \mathbf{h}_j) - \mathbf{w}_i \mathbf{h}_j] \\ & \text{subject to} && \mathbf{w}_{ik} \geq 0 \ \& \ \mathbf{h}_{kj} \geq 0 \end{aligned}$$

Algorithm Updates:

$$\begin{aligned} \hat{\mathbf{w}}_{ik} &= \hat{\mathbf{w}}_{ik} \left(\frac{\sum_{j=1}^p \hat{\mathbf{h}}_{kj} \mathbf{x}_{ij} / \hat{\mathbf{w}}_i^T \hat{\mathbf{h}}_j}{\sum_{j=1}^p \hat{\mathbf{h}}_{kj}} \right) \\ \hat{\mathbf{h}}_{kj} &= \hat{\mathbf{h}}_{kj} \left(\frac{\sum_{i=1}^n \hat{\mathbf{w}}_{ik} \mathbf{x}_{ij} / \hat{\mathbf{w}}_i^T \hat{\mathbf{h}}_j}{\sum_{i=1}^n \hat{\mathbf{w}}_{ik}} \right) \end{aligned}$$


Local solution.

NMF - Uses

- ① Dimension Reduction / Pattern Recognition.
 - ▶ Similar to PCA (e.g. component scatterplots) except that patterns of correlation found in the same direction.
- ② Archetypal Analysis.
 - ▶ Caricatures (segments; contrastive categorization) vs. Prototypes (averages).
- ③ Soft-clustering.
 - ▶ Discussed Next Lecture!

NMF - Archetypal Analysis


NMF



Original



\times




$=$



NMF - Archetypal Analysis

PCA



PCA vs. NMF

Similarities:

- Linear Dimension Reduction.
- Interpretation.

Differences:

- Factors are unordered.
- Factors NOT orthogonal.
- Changing K can fundamentally change factors.
- Non-unique, non-global solution.
- Depends on initialization. (Run several times and take the best).

Choosing K

Choice depends on goal:

- Dimension Reduction:
 - ▶ Residual sums of squares (or dispersion) - Screeplot.
- Clustering:
 - ▶ Consensus, silhouette, etc. (Discussed next lecture!).
- Archetypal Analysis:
 - ▶ Sparsity, factor purity, etc.

NMF - Summary

Strengths:

- Interpretation (often more appealing than PCA!).
- Applications - Clustering & Archetypal Analysis.
- Pattern Recognition.
- Others?

Weaknesses:

- Local solutions that depend strongly on K .
- Others?

In R: NMF package.

Independent Components Analysis (ICA)

Pre-processing Step: Reduce $\mathbf{X}_{n \times p}$ to $\tilde{\mathbf{X}}_{K \times p}$ with $K < n$ # independent sources. (Typically via PCA!)

Idea: $\tilde{\mathbf{X}}_{K \times p} \approx \mathbf{A}_{K \times K} \mathbf{S}_{K \times p}$.

- Assumption: $\tilde{\mathbf{X}}$ a matrix of K scrambled independent signals.
- $\mathbf{A}_{K \times K}$ *Mixing Matrix* - denotes how signals are scrambled to form sources in data.
- $\mathbf{S}_{K \times p}$ *Signal Matrix* - each row of \mathbf{S} is an independent signal.

PCA finds uncorrelated, but not independent signals.

ICA Uses

① Blind Source Separation.

- ▶ Assume K independent signals got scrambled, but record K scrambled versions of the signal.
- ▶ Cocktail Party Problem.


`http:
//research.ics.aalto.fi/ica/cocktail/cocktail_en.cgi`

② Denoising.

- ▶ Noise - independent from true signals.

ICA vs. PCA

Blind Source Separation:



ICA Algorithms

Fast ICA:

- Finds rotations of \mathbf{X} that are “non-Gaussian”.
- Uses non-Gaussian contrast functions:
 - ▶ $g(x) = x^4$.
 - ▶ $g(x) = \tanh(x)$.
- Generalization of projection pursuit.

Others:

- Infomax (entropy).

Not Statistically Independent!

PCA vs. ICA

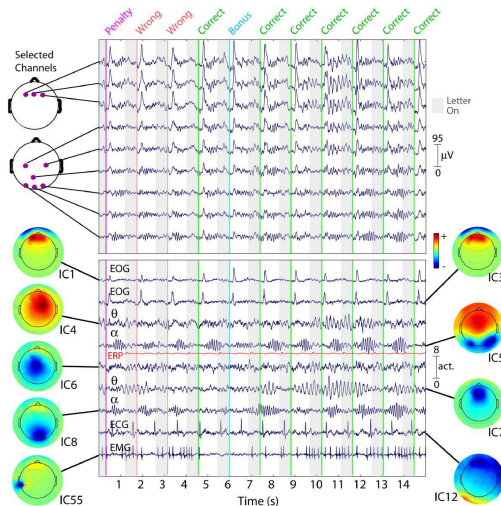
Similarities:

- Linear Dimension Reduction.
- Interpretation.

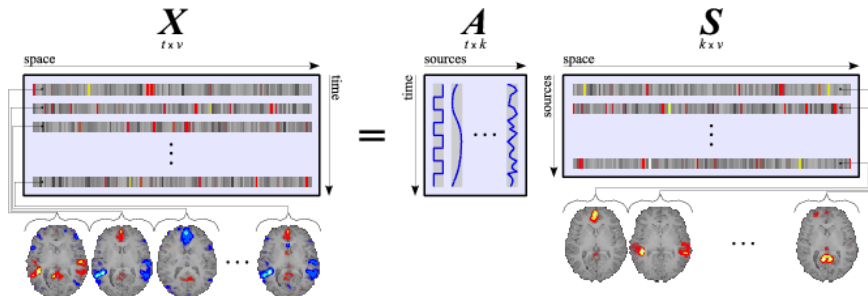
Differences:

- Factors are unordered.
- Factors NOT invariant - same solution by applying a permutation.
- Factors NOT orthogonal.
- Changing K can fundamentally change factors.
- Non-unique.
- No optimization criterion to evaluate solution.

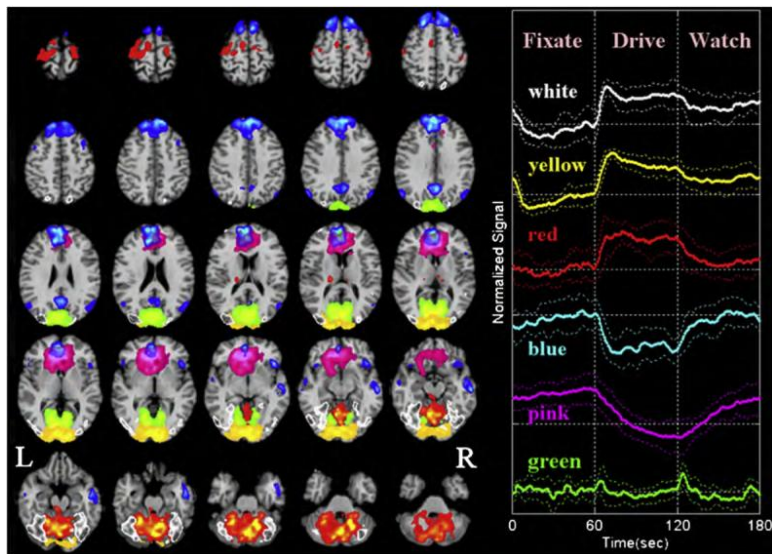
ICA Applications - EEG



ICA Applications - fMRI



ICA Applications - fMRI



ICA Summary

Strengths:

- Interpretation.
- Applications - Blind Source Separation & Denoising.
- Others?

Weaknesses:

- Solutions that depend strongly on K .
- Solutions can be rotated.
- Others?

In R: `fastICA` package.

Multidimensional Scaling (MDS)

Multidimensional Scaling (MDS)

Idea:

- Visually represent proximities (similarities or distances) between objects in a lower dimensional space.
- Input: Matrix of similarities or dissimilarities, $\mathbf{D}_{n \times n}$ (don't need the data itself!).
- Goal: Find projections ($\mathbf{z}_1, \dots, \mathbf{z}_K$ where $\mathbf{z} \in \mathbb{R}^n$) that preserve original distances in \mathbf{D} in a lower dimensional space ($K \ll n$).
- Distances preserved by optimizing a *stress function*.
- Non-linear dimension reduction.


MDS - Example

Consider the distances between nine American cities:


| | BOS | CHI | DC | DEN | LA | MIA | NY | SEA | SF |
|-----|------|------|------|------|------|------|------|------|------|
| BOS | 0 | 963 | 429 | 1949 | 2979 | 1504 | 206 | 2976 | 3095 |
| CHI | 963 | 0 | 671 | 996 | 2054 | 1329 | 802 | 2013 | 2142 |
| DC | 429 | 671 | 0 | 1616 | 2631 | 1075 | 233 | 2684 | 2799 |
| DEN | 1949 | 996 | 1616 | 0 | 1059 | 2037 | 1771 | 1307 | 1235 |
| LA | 2979 | 2054 | 2631 | 1059 | 0 | 2687 | 2786 | 1131 | 379 |
| MIA | 1504 | 1329 | 1075 | 2037 | 2687 | 0 | 1308 | 3273 | 3053 |
| NY | 206 | 802 | 233 | 1771 | 2786 | 1308 | 0 | 2815 | 2934 |
| SEA | 2976 | 2013 | 2684 | 1307 | 1131 | 3273 | 2815 | 0 | 808 |
| SF | 3095 | 2142 | 2799 | 1235 | 379 | 3053 | 2934 | 808 | 0 |

Can we represent these cities in a 2D space like a map?

MDS - Example



MDS - Example



MDS - Example



MDS - Stress Functions

- Input: $\mathbf{D}_{n \times n}$: $d_{ii'}$ denotes distance between object i and i' .
- Output: Projections, $\mathbf{z}_1, \dots, \mathbf{z}_k$, $\mathbf{z}_k \in \mathbb{R}^n$, that preserve distances.

Stress Functions:

- Least squares or Kruskal-Shephard Scaling:

$$S_D(\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_K) = \sqrt{\sum_{i \neq i'} (d_{ii'} - \|\mathbf{z}_i - \mathbf{z}_{i'}\|)^2}.$$



- Sammon mapping: preserve smaller pairwise distances

$$\sum_{i \neq i'} \frac{(d_{ii'} - \|\mathbf{z}_i - \mathbf{z}_{i'}\|)^2}{d_{ii'}}.$$

- Shepard-Kruskal nonmetric scaling ($\theta(\cdot)$: an increasing function):

$$\frac{\sum_{ii'} [\theta(\|\mathbf{z}_i - \mathbf{z}_{i'}\|) - d_{ii'}]^2}{\sum_{ii'} d_{ii'}^2}.$$

MDS - Example



MDS Properties

- Data not needed - only dissimilarities.
- Algorithm - gradient descent.
- Choosing K :
 - ▶ Scree plot (like PCA).
 - ▶ Shepard Diagram - plot proximities against distances in Z .
- Interpreting MDS maps:
 - ▶ Axes and orientation arbitrary.
 - ▶ Can be rotated.
 - ▶ Only relative locations important.
 - ▶ Typically looks for objects close in the MDS map.

MDS vs. PCA

Similarities:

- Dimension reduction for visualization.

Differences:

- Non-linear vs. Linear.
- Local solution & arbitrary map.
- Non-unique & local solution.

MDS - Summary

Strengths:

- Visualizing proximities.
- Only need dissimilarities.
- Others?

Weaknesses:

- Arbitrary maps.
- Which stress function?
- High-dimensional settings? ($p \gg n$ - more features than objects)
- Others?

In R: `dist`; `cmdscale` - classical MDS; `isoMDS` - Kruskals's MDS and `sammon` in MASS package.

Dimension Reduction Wrap-Up

Techniques Covered:

- PCA.
- NMF.
- ICA.
- MDS.

Dimension Reduction Wrap-Up

Comparative Strengths & Weaknesses:

| Property | PCA | NMF | ICA | MDS |
|----------|-----|-----|-----|-----|
| ⋮ | ⋮ | ⋮ | ⋮ | ⋮ |

References

Textbooks:

- Elements of Statistical Learning by Hastie, Tibshirani & Friedman.
<http://statweb.stanford.edu/~tibs/ElemStatLearn/>

Some of the figures in this presentation are taken from this textbook with permission from the authors.